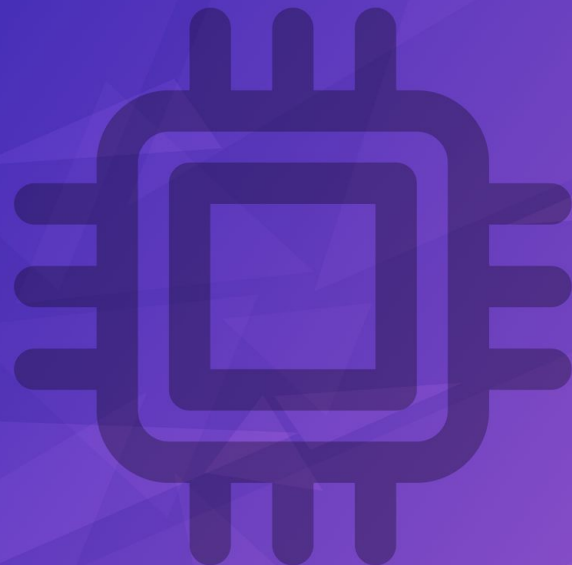


MACHINE LEARNING IN PRACTICE WITH SPARK MLlib: AN INTELLIGENT DATA ANALYZER.

STRATA+HADOOP WORLD

Eiti Kimura / Flávio Clésio
@Movable Brazil - Dec 2016



ABOUT US



Flávio Clésio

- Core *Machine Learning* at Movable
- MSc. in Production Engineering (Machine Learning in Credit Derivatives/NPL)
- Specialist in Database Engineering and Business Intelligence
- Blogger at *Mineração de Dados* (Data Mining) - <http://mineracaodedados.wordpress.com>



flavioclesio

ABOUT US



Eiti Kimura

- Software Architect and TI Coordinator at Movable
- Msc. in Electrical Engineering
- Apache Cassandra MVP (2014/2015 e 2015/2016)
- Apache Cassandra *Contributor* (2015)
- Cassandra Summit *Speaker* (2014 e 2015)
- Cassandra Summit *Reviewer* (2016)



eitikimura

WE MAKE LIFE BETTER
THROUGH OUR APPS

mobile

Mobile is the company behind of several
apps that makes the life easier



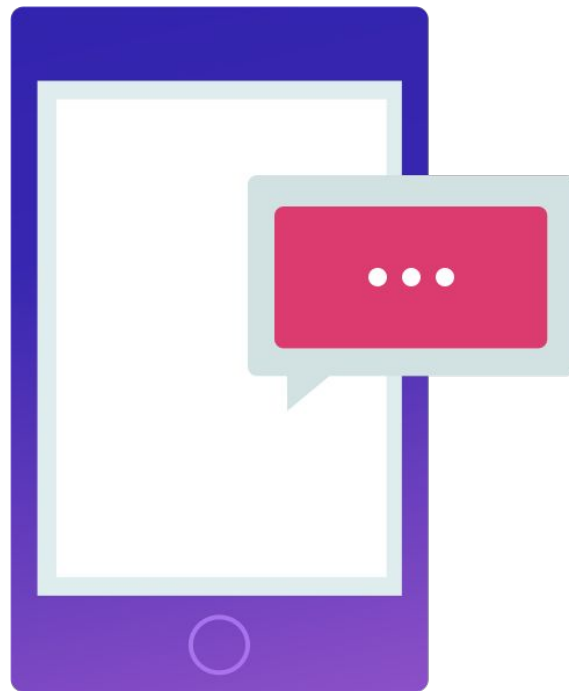
THE BEST CONTENT FOR KIDS

PlayKids is the #3 Top Grossing worldwide App for Children's. Kid Safe Toddler App.



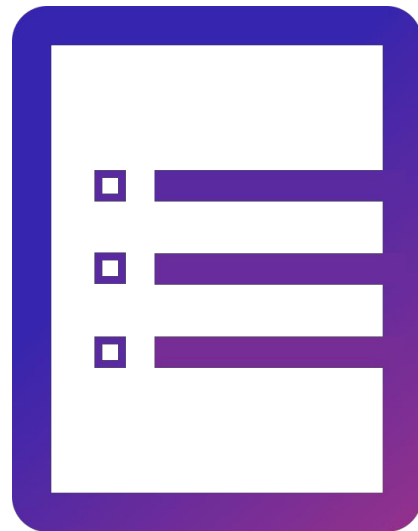
MESSAGING AND BILLING SERVICES FOR MOBILE CARRIERS

Huge case of SMS services for corporative
and mobile content distribution.



AGENDA

- The Movie's Platform Case
- Regression Modeling: a bit of theory
- Practical Machine Learning Model Training
- Presenting Watcher-ai
- Results and Conclusions




MOBILE'S SUBSCRIPTION AND BILLING PLATFORM IN ITS SIMPLEST FORM

- A distributed platform
- User's subscription management
- MISSION CRITICAL platform: can not stop under any circumstance



MAIN PROBLEM: MONITORING

How can we check if platform is fully functional based on data analysis only?



Tip: what if we ask help to an intelligent system?



HOW DATA LOOKS LIKE?

- 120 Millions + of billing requests attempt a day
- 4 main mobile carriers drive the operational work

carrier weight	date time	avg resp. time	succ. charges	no credit	general errors	total attempts
1	2016-10-31 0-8 pm	1014 ms	99.107	24.232.849	3.239.499	27.571.455
1	2016-11-01 0-8 pm	1204 ms	106.232	23.989.076	4.024.136	28.119.444
1	2016-11-02 0-8 pm	1186 ms	114.013	24.513.752	3.217.619	27.845.384
1	2016-11-03 0-8 pm	1117 ms	118.110	23.714.608	3.205.513	27.038.231
1	2016-11-04 0-8 pm	1138 ms	124.246	22.553.776	5.135.307	27.813.329
1	2016-11-05 0-8 pm	942 ms	102.674	23.556.432	4.072.168	27.731.274

DATA AND ALGORITHM MODELING

Sample of data (predicting the number of success)

The diagram illustrates the structure of a dataset. It shows a sequence of values: 61.083, [4.0, 17h, 3.0, 1259.0, 24.751.650, 2.193.67, 26.314.551]. A bracket below the first value (61.083) is labeled 'label/target'. A larger bracket below the remaining values in the list is labeled 'features'.

# success	carrier_weight	hour	week	response_time	#no_credit	#errors	# attempts
61.083,	[4.0,	17h,	3.0,	1259.0,	24.751.650,	2.193.67,	26.314.551]

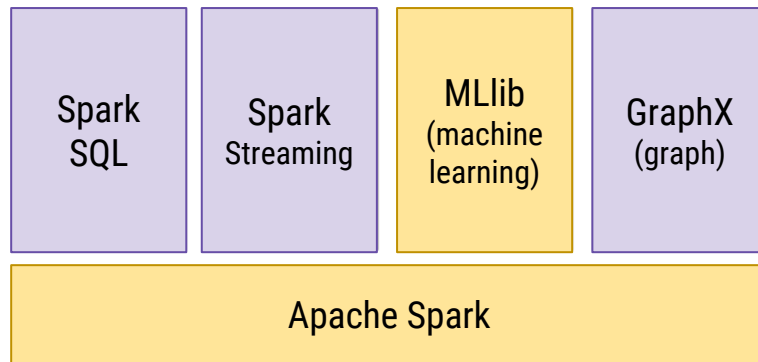
label/target features

SUPERVISED LEARNING

Linear Regression

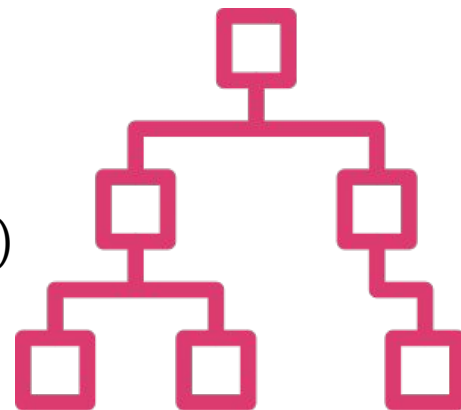
SPARK MLLIB

- MLlib is Apache Spark's scalable machine learning library.
- MLlib contains many algorithms and utilities, including Classification, Regression, Clustering, Recommendation, Pipelines and so on...



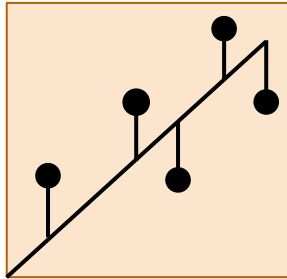
TESTED ALGORITHMS USING SPARK ML LIB

- Linear Model with Stochastic Gradient (SDG)
- Lasso with SGD Model (L1 Regularization)
- Ridge Regression with SGD Model (L2 Regularization)



REGRESSION MODELING: A BIT OF THEORY

- Linear regression it's a statistical method that investigates the relationship and interdependency between variables to get a numerical result.



Regression Algorithms

A LITTLE BIT OF MATH

$$y = \alpha + (\beta_1 * x_1) + (\beta_2 * x_2) + (\beta_n * x_n) + \epsilon$$

as:

y = Value to be predicted (dependent variable)

α = Intercept (Where the slope gets the Y axis and the value of X is 0) - Endogenous Factors

β = Regression Coefficients

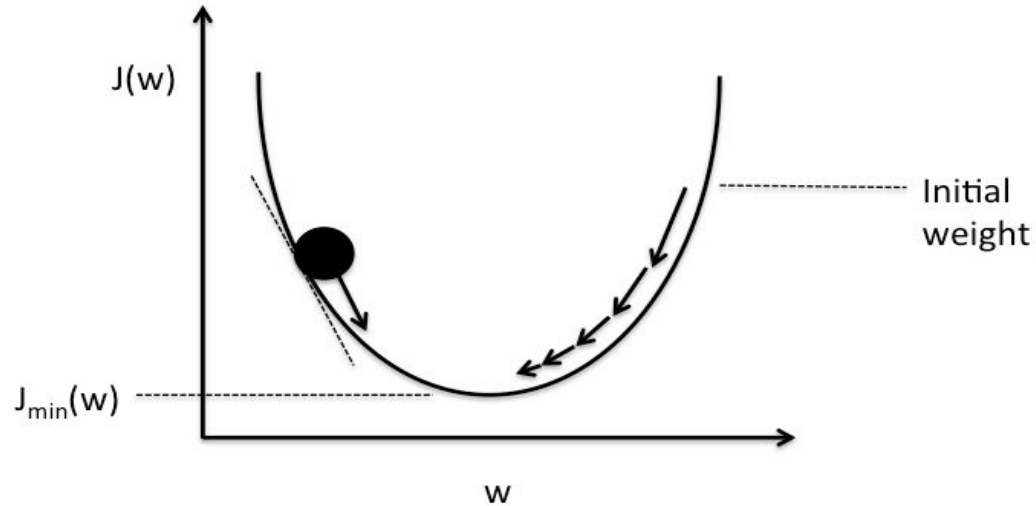
$x_1 \dots x_n$ = Values of independent variables (e.g. columns of a database)

ϵ = Random noise, non explicit errors - Exogenous Factors

LET'S TALK ABOUT REGULARIZATION

- To avoid the overfitting problem **Spark MLlib** embed some regularization methods like LASSO (L1) and Ridge (L2).
- LASSO (L1) regularization have as feature a penalty adds the penalty equivalent to the absolute value of the magnitude of the coefficients.
- Ridge L2) regularization method the penalty is equivalent to the magnitude of the coefficients raised to the square.

STOCHASTIC GRADIENT DESCENT



Schematic of gradient descent.



THINK ABOUT A TAILOR MADE SUIT

STOCHASTIC GRADIENT DESCENT

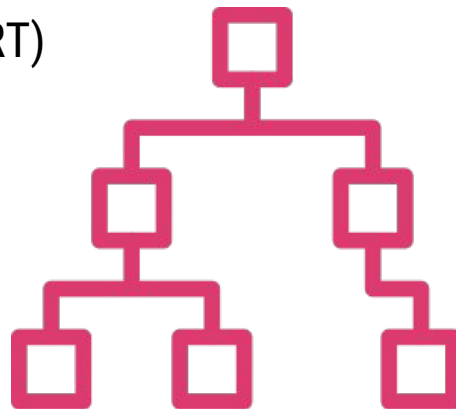


- Example: Stochastic Gradient Descent ([Tailoring a Suit, from Quora](#))

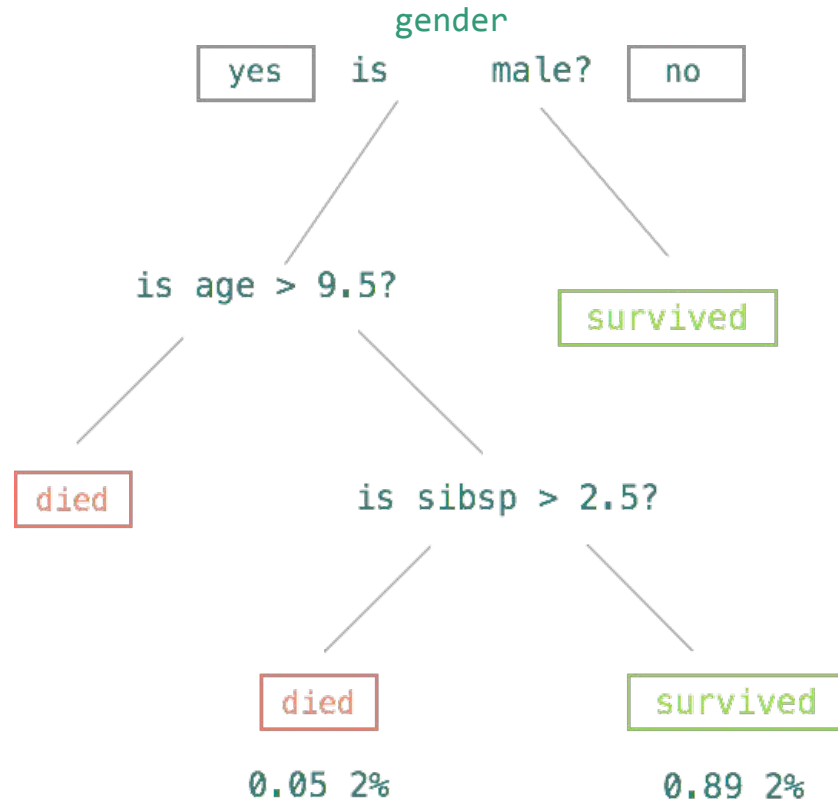
- 1) Tailor makes initial estimate. *(See the parameters of the model)*
- 2) A random guy (or a subset of the full group) tries the **suit** and gives feedback.
(take a sample of dataset)
- 3) Make a small adjustment according to feedback. *(change to reduce the error of the model)*
- 4) While the tailor still have time for this, he goes to 2. *(iterate and repeat the process)*

TESTED ALGORITHMS

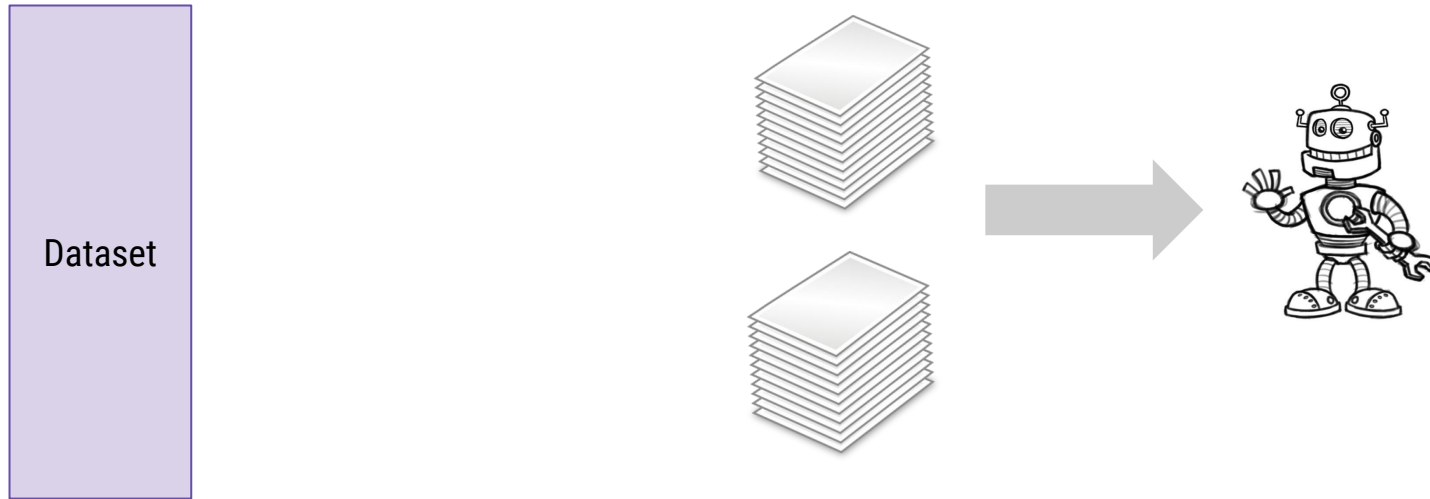
- Decision Tree with Regression Properties (Not CART)



REGRESSION TREE



MODELING LIFECYCLE STEP-BY-STEP



SPARK NOTEBOOK

<http://spark-notebook.io/>



SPARK NOTEBOOK

Regression-Benchmarks

Scala [2.10.4] Spark [1.6.0] Hadoop [2.2.0] {Hive ✓} {Parquet ✓} ○

File Edit View Insert Cell Kernel Help



Code



Cell Toolbar: None



```
In [ ]: def buildLabelValue(list: List[Double]) : Double = {  
  // index = 4 is the number of success of the hour,  
  return if (list(4) != 0.0) Math.log(list(4)) else  
}  
  
def buildFeatures(list: List[Double]) : List[Double]  
  // remove the index 4, which means the number of  
  return list.patch(4, Nil, 1)  
}
```

```
In [ ]: // reading pre processed dataset  
val rdd = sc.objectFile[List[Double]](ROOT_DIR + "/rc  
  
// building the LabelPoint, using success as Label  
val labelSet = rdd.map{l => val label = buildLabelVe  
                           val features = buildFeat  
                           LabeledPoint(label, Vect  
rdd.take(10).foreach(println)
```

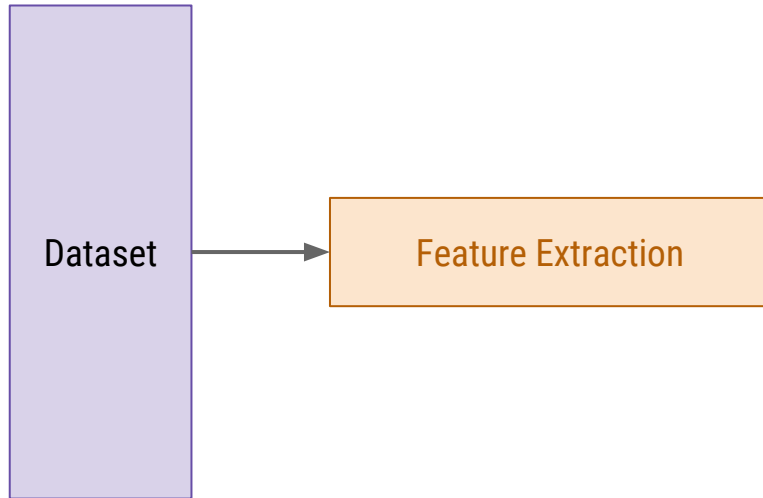
LOADING DATA

scala code snippet

```
// reading dataset  
val rdd = sc.objectFile[List[Double]](ROOT_DIR + "/rdd-processed")
```

carrier_weight	hour	week	resp_time	# success	no_credit	# errors	# attempts
List(4.0,	17.0,	3.0,	1709.4,	39511.8,	2386316.3,	291279.6,	2717107.8)
List(2.0,	8.0,	5.0,	749.9,	51910.5,	1.27E7,	1951005.1,	1.47E7)
List(4.0,	11.0,	5.0,	1690.0,	18519.0,	562289.5,	173717.3,	754525.9)
List(2.0,	22.0,	1.0,	911.4,	257598.2,	4.05E7,	1.3E7,	5.4E7)
List(4.0,	7.0,	5.0,	1386.3,	1775.3,	391668.5,	75062.6,	468506.5)
List(1.0,	23.0,	4.0,	561.8,	195032.6,	2.8E7,	5279717.1,	3.41E7) ...

MODELING LIFECYCLE




FEATURE EXTRACTION

scala code snippet

```
def buildLabelValue(list: List[Double]) : Double = {  
  // index = 4 is the number of success, that is what we want to predict  
  return if (list(4) != 0.0) Math.log(list(4)) else 0.0  
}  
  
def buildFeatures(list: List[Double]) : List[Double] = {  
  // remove the index 4, which means the number of success  
  return list.patch(4, Nil, 1)  
}  
  
// building the LabelPoint, using success as Label  
val labelSet = rdd.map{l => val label = buildLabelValue(l)  
                           val features = buildFeatures(l)  
                           LabeledPoint(label, Vectors.dense(features.toArray))}
```

labelSet: RDD[org.apache.spark.mllib.regression.LabeledPoint]

MODELING LIFECYCLE



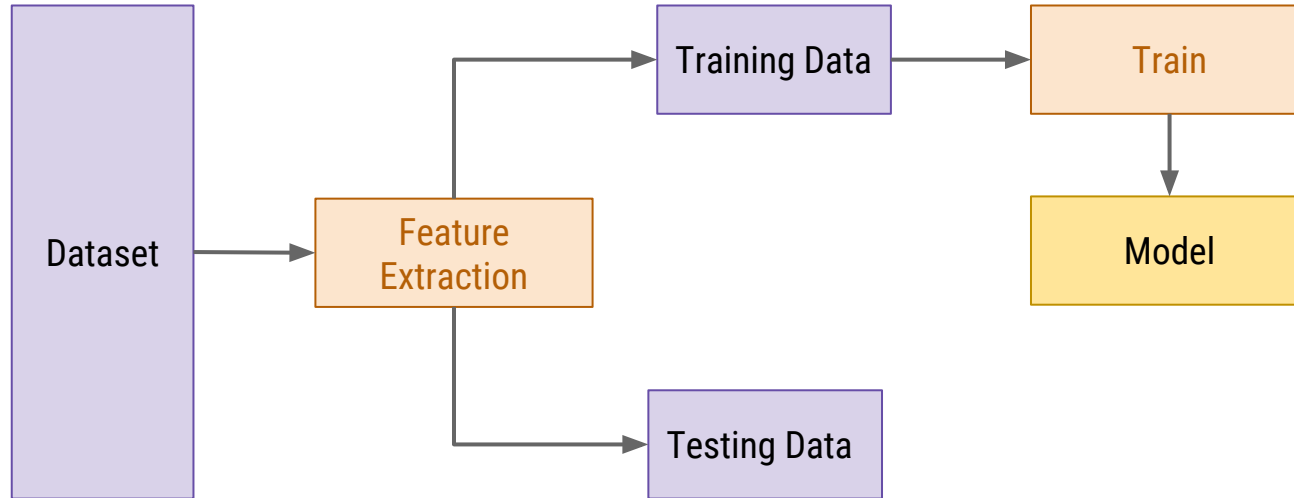
SPLITTING DATASET

scala code snippet

```
//Split data into training and test
val splits = labelSet.randomSplit(Array(0.70, 0.30), seed = 13L)
val training = splits(0)
val test = splits(1)
```

The main idea is to use 70% of data to train the model and 30% to evaluate the model performance

MODELING LIFECYCLE





standardizes

features by scaling



filtering

the dataset and grouping

AUXILIARY FUNCTIONS

scala code snippet

```
// 1 - standardizeTrainingSet
val scaler = new StandardScaler(withMean = false, withStd = true)
                .fit(rdd.map(x => x.features))
// carrier weight list for filtering
val range = List(1.0, 2.0, 4.0, 5.0)
```

```
// 2 - filterTrainingSet
range.map{carrierWeight =>
    val trainingSet = rdd.filter(l => l.features.apply(0) == carrierWeight)
                          .map(x => LabeledPoint(x.label,
                                                  scaler.transform(x.features)))
    (idx, trainingSet)
}.toMap
```



Linear Regression SGD

Ridge Regression SGD

Lasso Regression SGD

Decision Tree w/ Regression



the champion is...

TRAINING: LINEAR REGRESSION WITH SGD

scala code snippet

```
def buildSGDModelMap(rdd:Map[Double, RDD[LabeledPoint]]) : Map[Double,
LinearRegressionModel] = {

  val carrierWeight = List(1.0, 2.0, 4.0, 5.0)
  return carrierWeight.map{idx =>
    // Building the model
    val numIterations = 100
    var regression = new LinearRegressionWithSGD().setIntercept(true)
    regression.optimizer.setStepSize(0.1)
    regression.optimizer.setNumIterations(numIterations)

    // get dataset for a specific carrier weight
    val dataset = rdd.get(idx).orElseNull;
    (idx, regression.run(dataset)) //<< training starts here
  }.toMap
}

val mapTraining = standardizeTrainingSet(training)
val mapSGDModel = buildSGDModelMap(mapTraining) //map with a model for each carrier
```

TRAINING: LASSO/RIDGE REGRESSION

scala code snippet

```
// instantiating the algorithm and setting the params
val regression = new LassoWithSGD()
regression.optimizer.setStepSize(0.1)
regression.optimizer.setNumIterations(100)

// training model
val model:LassoModel = regression.run(dataset)
```

```
// instantiating the algorithm and setting the params
val regression = new RidgeRegressionWithSGD()
regression.optimizer.setStepSize(0.1)
regression.optimizer.setNumIterations(100)

val model:RidgeModel = regression.run(dataset)
```

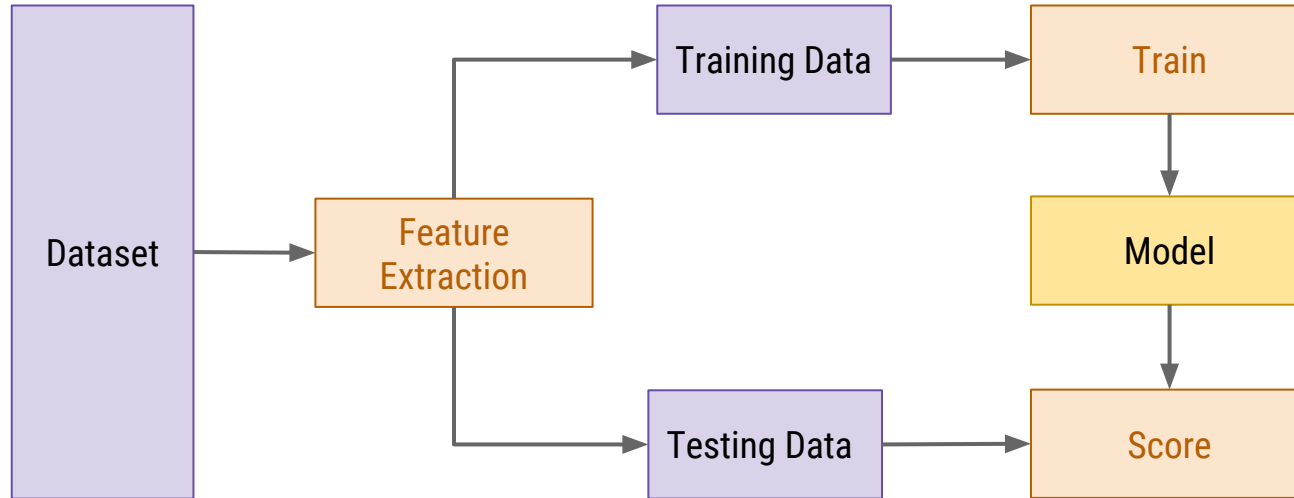
TRAINING: DECISION TREE WITH REGRESSION

scala code snippet

```
val categoricalFeaturesInfo = Map[Int, Int]()
val impurity = "variance" // indicates it is a regression tree
val maxDepth = 7 // the tree's max depth
val maxBins = 32 // max number of bins (groups)

val model = DecisionTree.trainRegressor(filteredSet, categoricalFeaturesInfo,
                                         impurity, maxDepth, maxBins);
```

MODELING LIFECYCLE



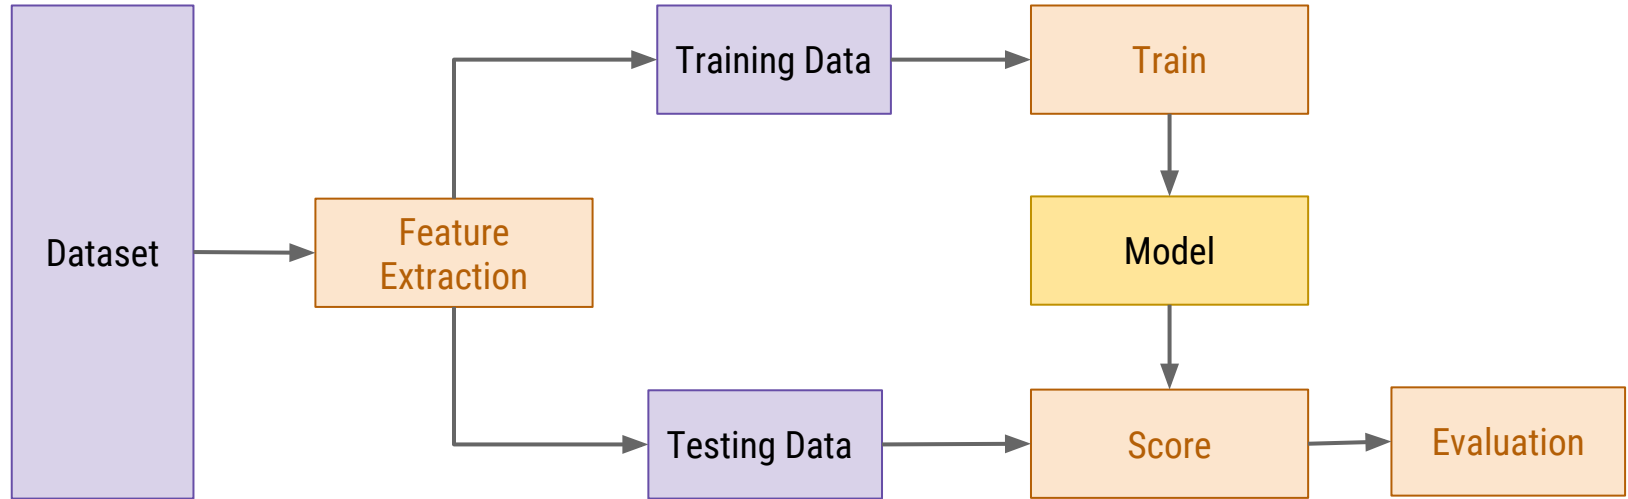
SCORE: AUXILIARY FUNCTIONS

scala code snippet

```
// predict values using the trained model
val labelsAndPredictions = test.map { point =>
  val carrier = point.features.apply(0) // 1st param is the carrier weight
  val model = getModelForCarrier(carrier)
  val prediction = model.predict(point.features) //<< prediction happens here
  (point.label, prediction)
}
```

Expected Value	Predicted Value
(12.196509845132933,	9.97275651498185),
(11.956245114516188,	11.632408901614912),
(11.840353189256883,	10.02762098460309),
(12.130296102598983,	9.320716165463033),
(11.682180075417563,	11.503286266374285),
(12.094170705574166,	13.010471821918166),
(11.832490497556401,	6.910122430921404) ...

MODELING LIFECYCLE



MODEL EVALUATION CRITERIA

- OMTM (One Metric That Matters)
- Root Mean Squared Error (RMSE)


$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

TIME TO EVALUATE THE RESULTS

Machine Learning Tested Model	Accuracy	RMSE
Linear Model With SGD	44%	1.64
Lasso with SGD Model	43%	1.65
Ridge Regression with SGD Model	43%	1.66
Decision Tree Model	98%	0.20



WATCHER-AI INTRODUCTION



Hi I'm Watcher-ai!
It is nice to see you here



[watcher-ai] [warn] abnormal number ... ✕

of errors abnormal for carrier: 4 , current value: 24448.000000000015, predicted value: 12005.999999999995, prediction error: 103.63151757454628% , features: [4.0, 9.0, 4.0]. The number of errors are higher than

✕ Dispensar



Pushbullet API



Pre trained ML
models

Load

Query Data



Analytic Database

[watcher-ai] possible problem (succ error: 41.3%, carrier: 4)

of success charge error: 41.3 %critical,

of attempts error critical: 100.7 %

carrier: 4.0, hour: 11,

measured: 596557.0, predicted: 1197382.74

response time measured(ms): 991.0,

response time predicted(ms): 1780.9

**% of calculated
number of attempts
prediction error**


**informs % of error
and carrier weight**

**measured and
predicted values for
number of attempts
and response time**

MOVING TO *results*

PRELIMINARY RESULTS

- Was designed to be *Last barrier of defense* system
- Helped to catch troubles in the last 9 months
- Brought light in several problems of monitoring at Movile
- Catch any discrepancy in hourly fashion



IT'S JUST THE

beginning

It isn't the end.

IT IS JUST A WARM UP

- Automatic refeed and training using collected data. Analyse more data to predict possible errors with carrier
- Notify more people and specific teams (more complex problems)
- Refactory to be more generic, so other teams can add their own algorithm
- Why not interact with Watcher to guide the analysis?
- Don't limit your mind, there is a lot to keep improving...

THANK YOU!

STRATA+HADOOP WORLD



github.com/fclesio/watcher-ai-samples

@eitikimura



eiti-kimura-movile

eiti.kimura@movile.com

@flavioclesio



fclesio

flavio.clesio@movile.com